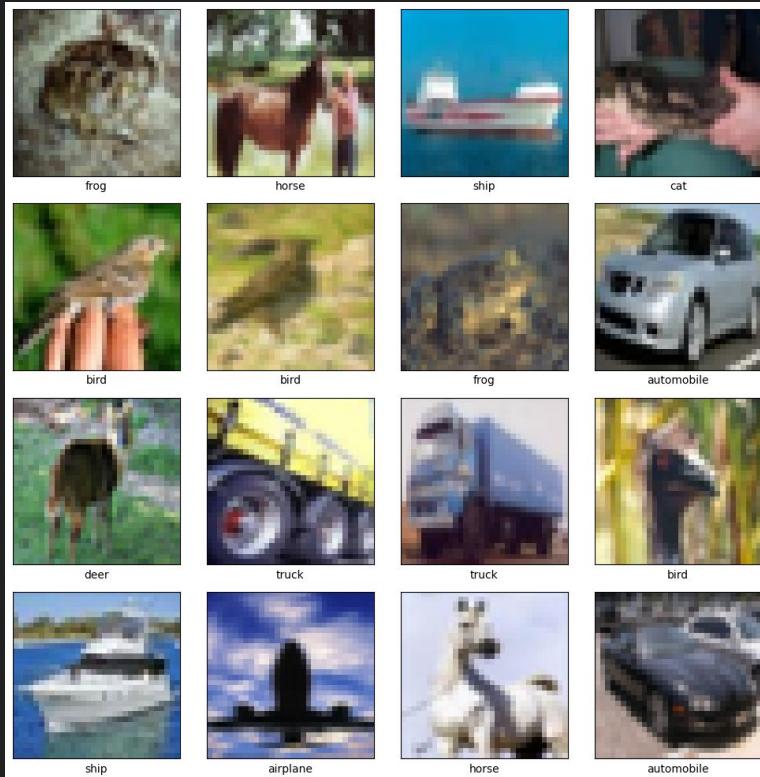




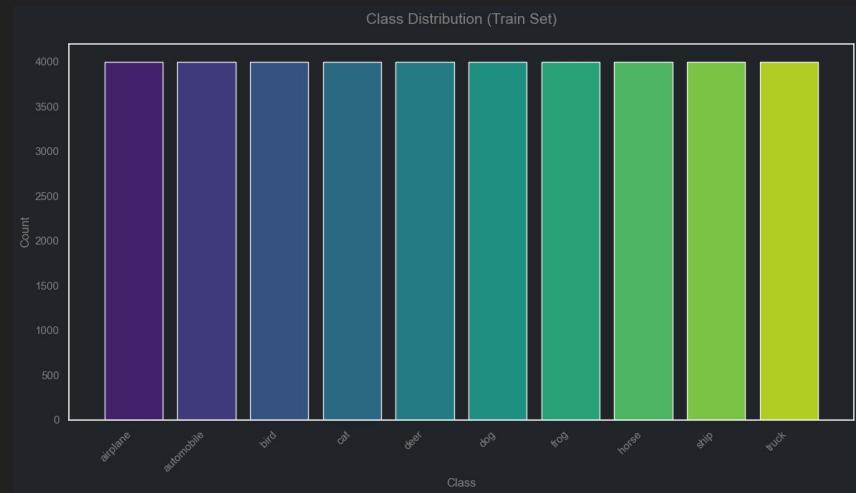
Computer vision

Alexandre, Janete, Isis
Feb 2026

CIFAR10 small images classification dataset



- [cifar10 dataset from keras](#)
- 60,000 32x32 colour images in 10 classes
 - training: 50,000 images
 - test: 10,000 images
- 6,000 images per class



Data Preparation

data splitting:

cifar10 is pre-split into 50K train and 10K test images, but we adjusted to

- 40,000 train
- 10,000 validation
- 10,000 test

Images

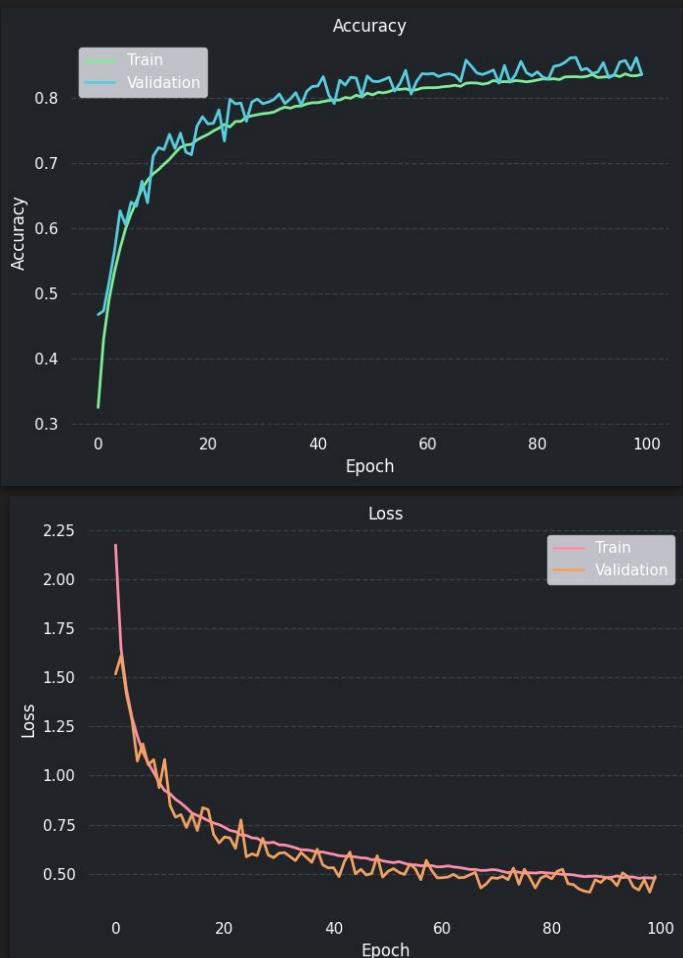
- Normalization

Target

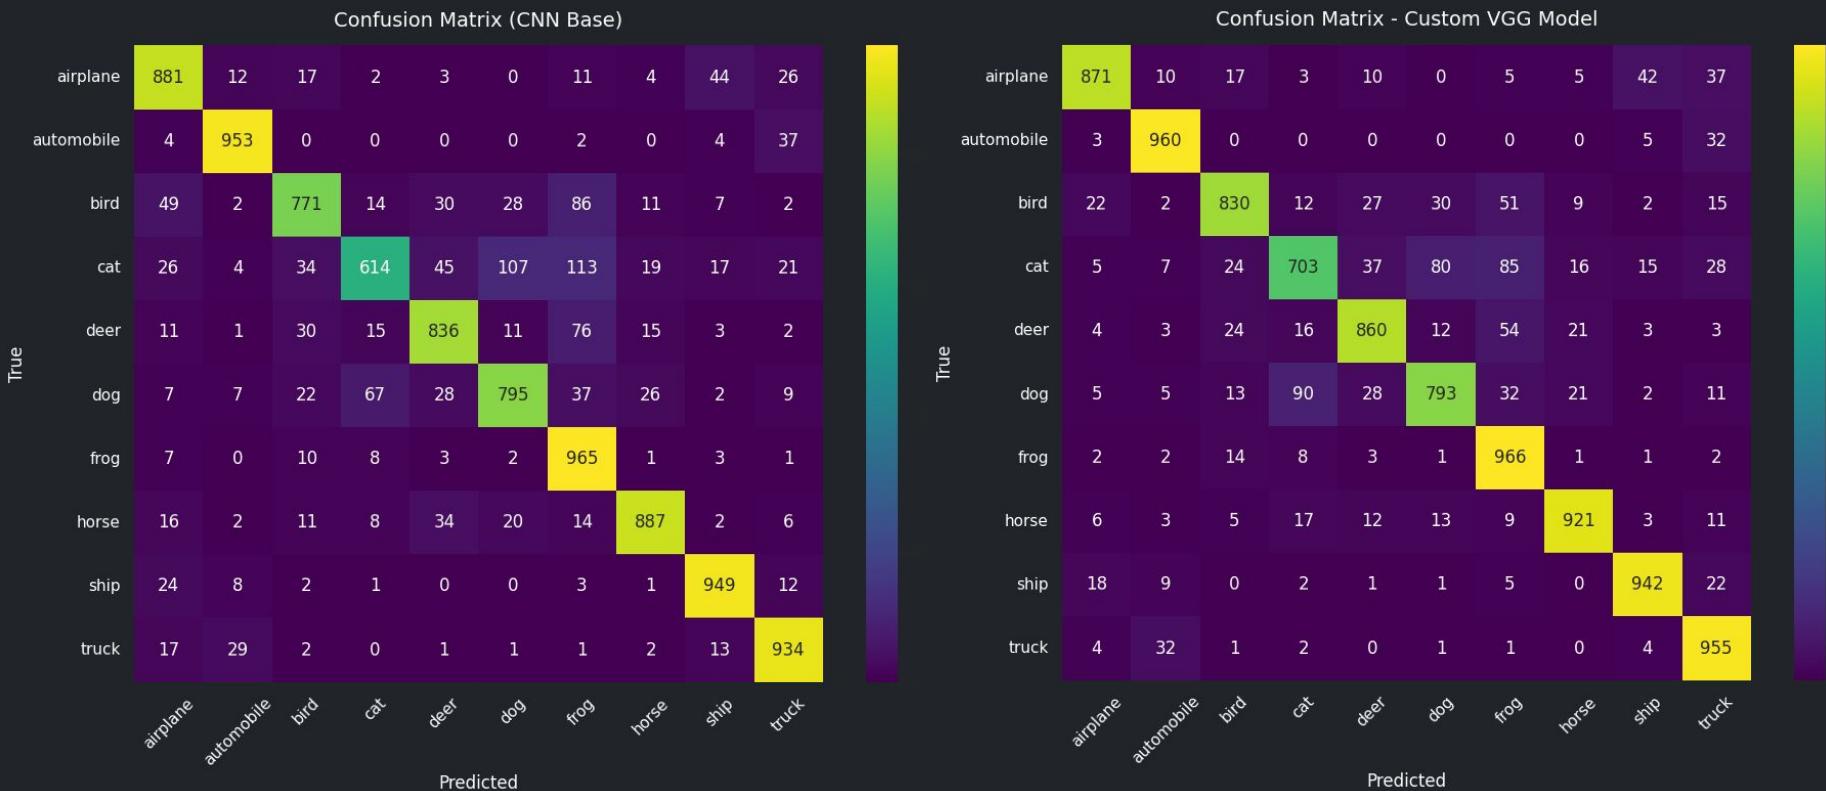
- One-Hot-Encoding

CNN – base model

- 3 blocks
 - Conv2D 32 (x2) ; 64 (x2) ; 128
 - activation “relu”
 - Dropouts (0.25,0.25,0.4)
- Classifier
 - Dense 256
 - Dropout 0.5
 - activation “softmax”
- Compiler:
 - adam, with adjusted learning rate
 - loss="sparse_categorical_crossentropy"
- Early Stop:
 - monitor val_loss
 - patience = 10
- ReduceLROnPlateau:
 - monitor val_loss
 - patience = 5



Confusion Matrixes CNN Models



Transfer Learning: EfficientNetB0 (04 Model)

EfficientNetB0



Adjustments:

- resized images to (96,96)
- randomly adjusted brightness, contrast, hue, saturation, random crop
- ‘unnormalised’ data, since image normalisation happens inside EfficientNet

EfficientNetB0 - Layers Unfrozen

Adjustments:

- same as before
- Unfroze half of the layers

Transferlearning: EfficientNetV2B1 (Model 06)

EfficientNetV2B1



Adjustments:

- resized images to (96,96)
- randomly adjusted brightness, contrast, hue, saturation, random crop
- ‘unnormalised’ data, since image normalisation happens inside EfficientNet

EfficientNetV2B1 – Layers Unfrozen

Adjustments:

- same as before
- Unfroze half of the layers

Transferlearning: EfficientNetV2S (Model 07)

PyTorch



Adjustments:

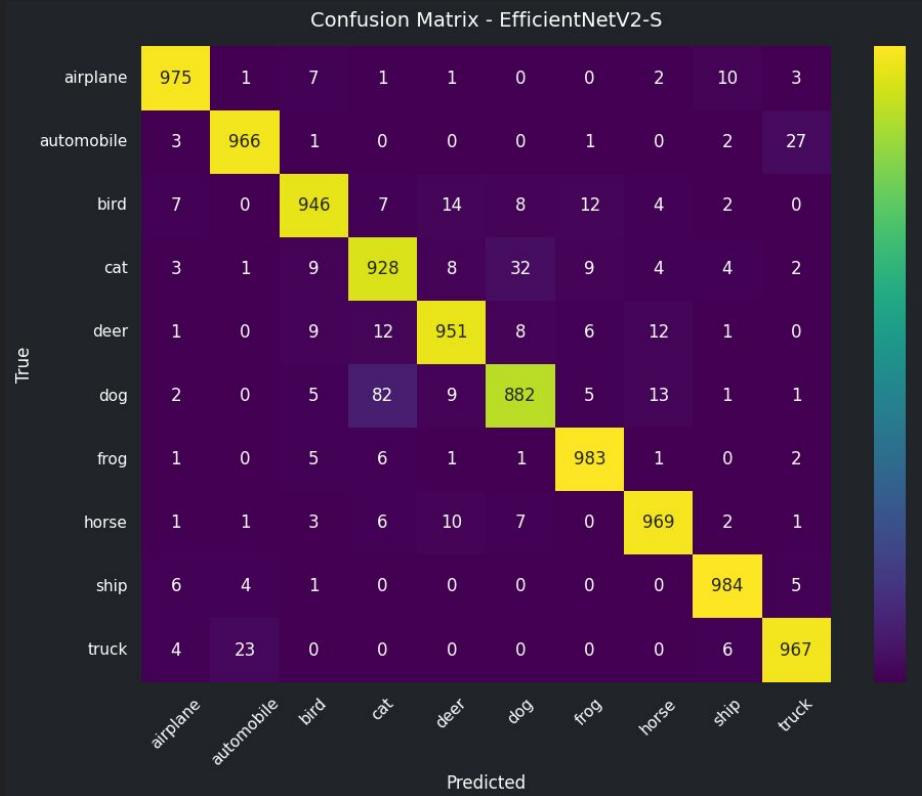
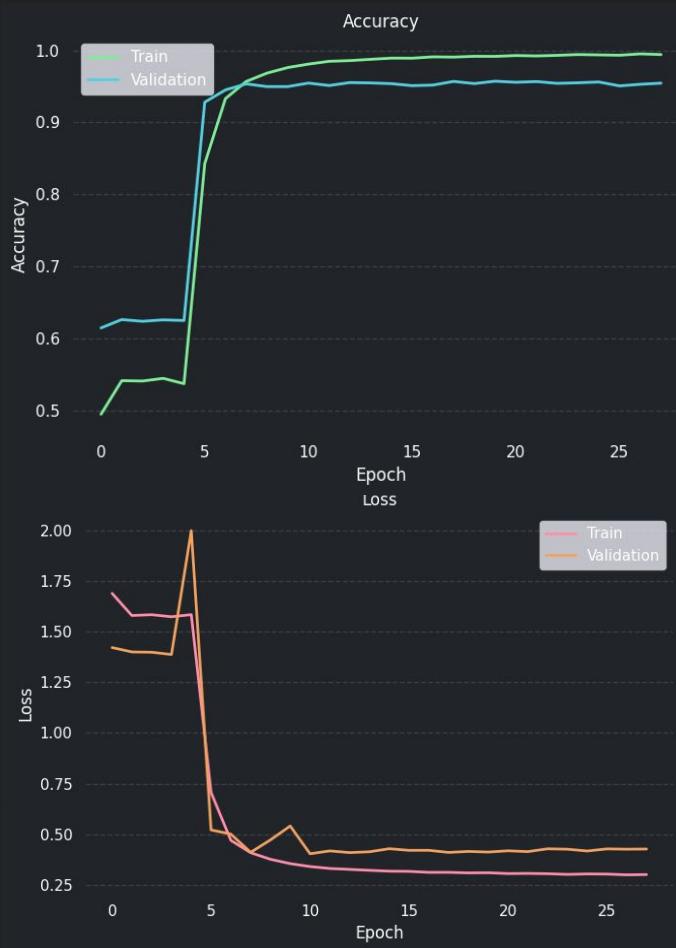
- using Pytorch library instead of keras
- same as 04 model
- Normalized by the mean and std

PyTorch - Layers Unfrozen

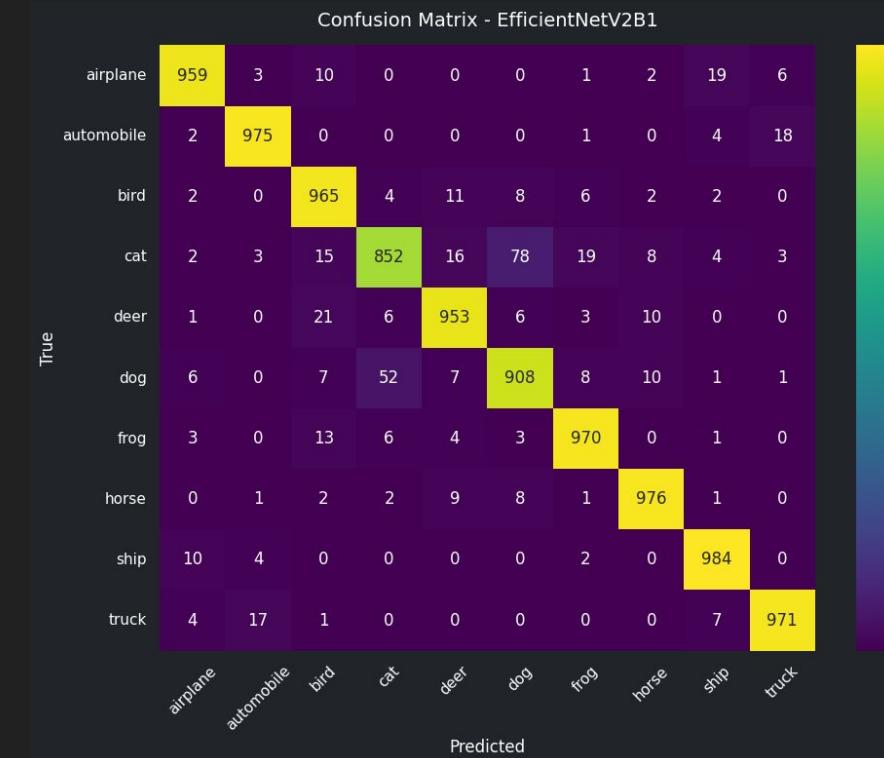
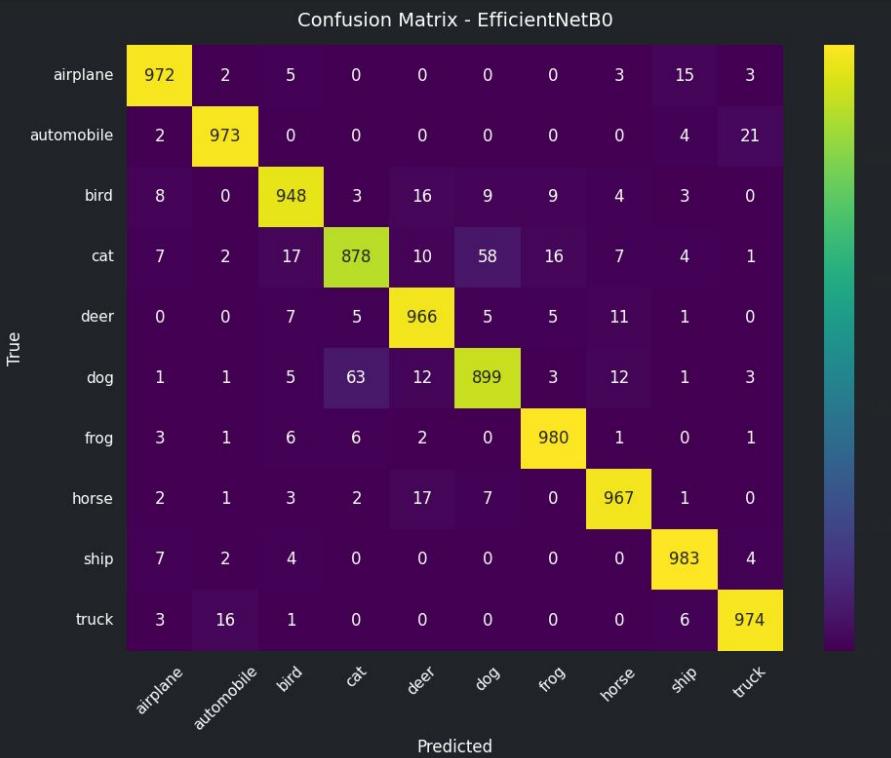
Adjustments:

- Unfroze half of the layers
- Smaller learning rate and label smoothing
- Increase the patience

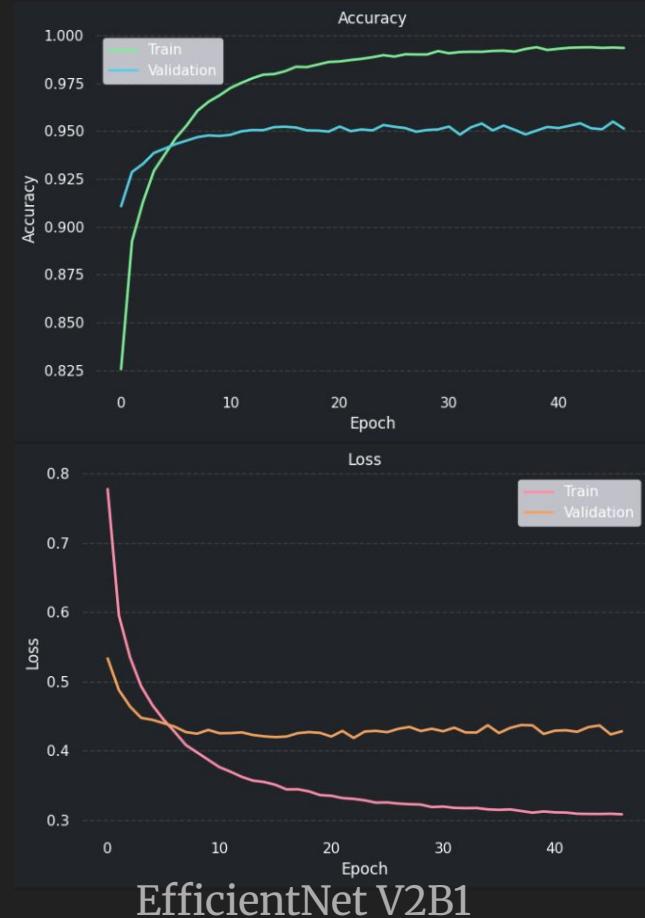
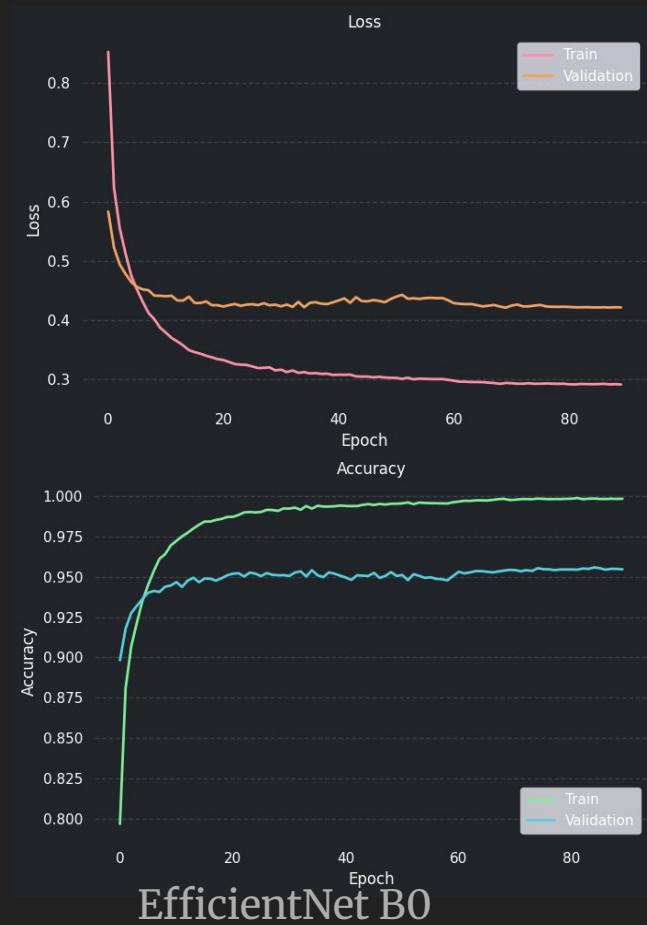
Best Model: PyTorch: Performance



Transfer Learning: Confusion Matrixes



Transfer Learning Curves



Overview: Adjustments made on...

Images:

- rotation
- adjusting saturation / hue / brightness
- cutting random holes into the images
- flipping
- random crop

Model Architecture:

- add layers
- batch normalisation
- replicating the architecture of the VGGBN model on our CNN model

Model Training

- learning rate scheduler:
 - reduces learning rate as you progress
- label smoothing:
 - accepts less confident predictions, improving performance

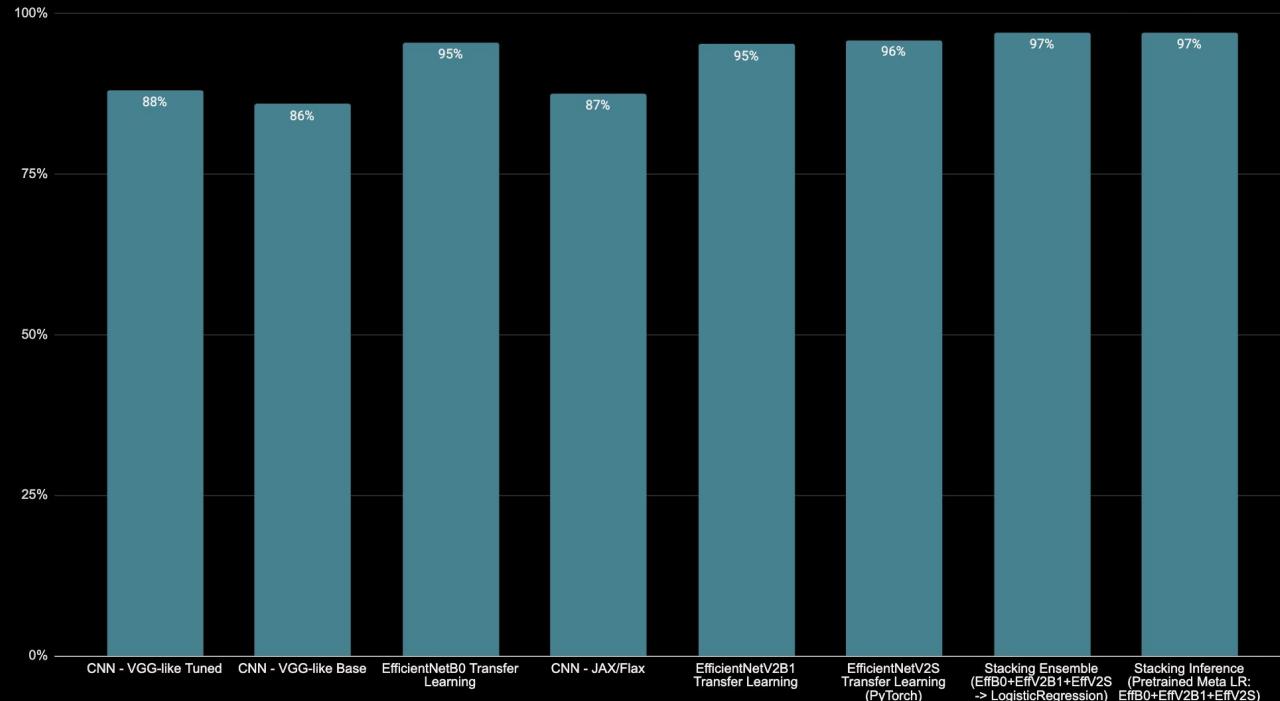
Very Best Model: Franceschina Model

- Stacking Models 04 + 06 + 07
- using a logistic regression to stack these models
- F1 Score: 96.94%



Evaluation Overview

Performance Summary: Accuracy



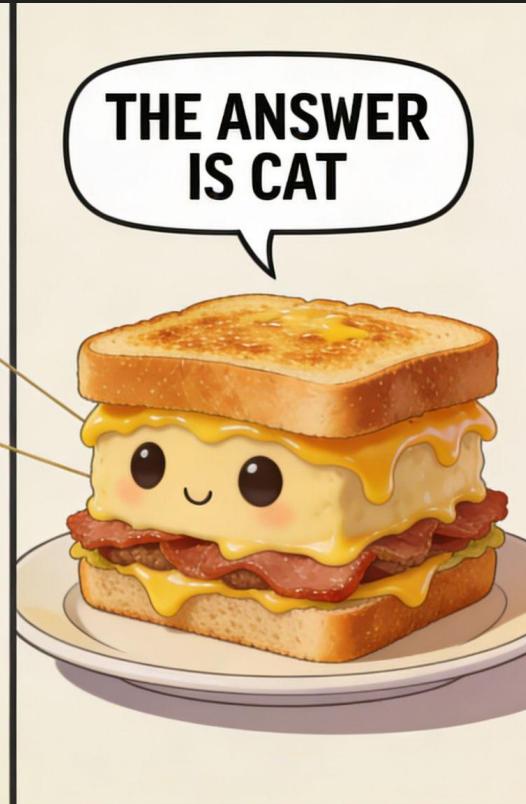
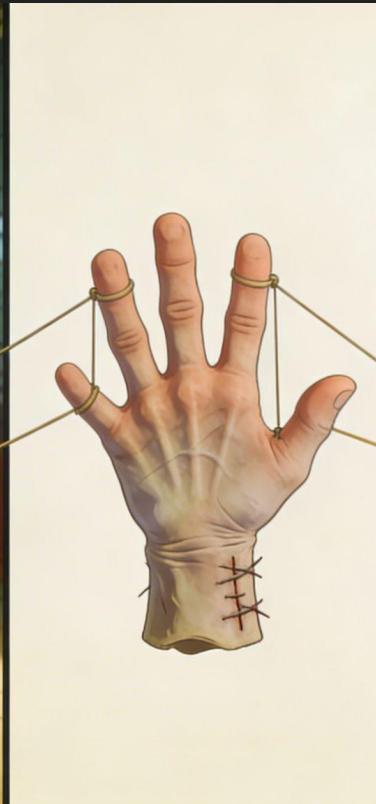
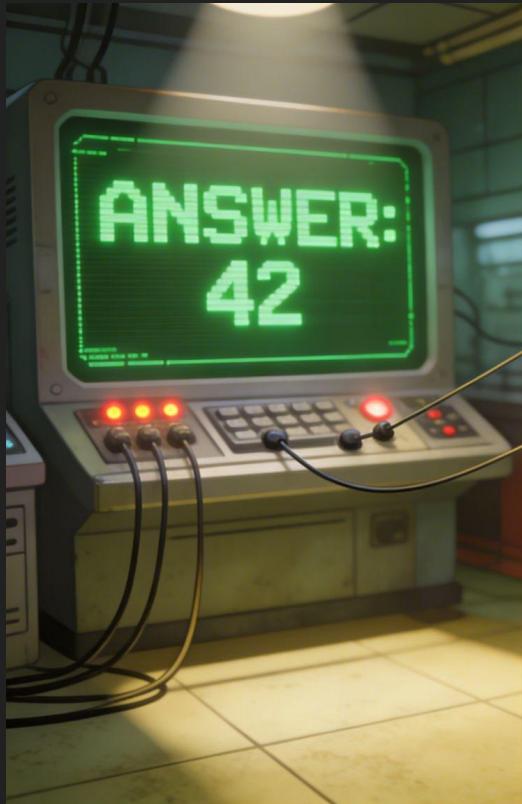
Conclusion

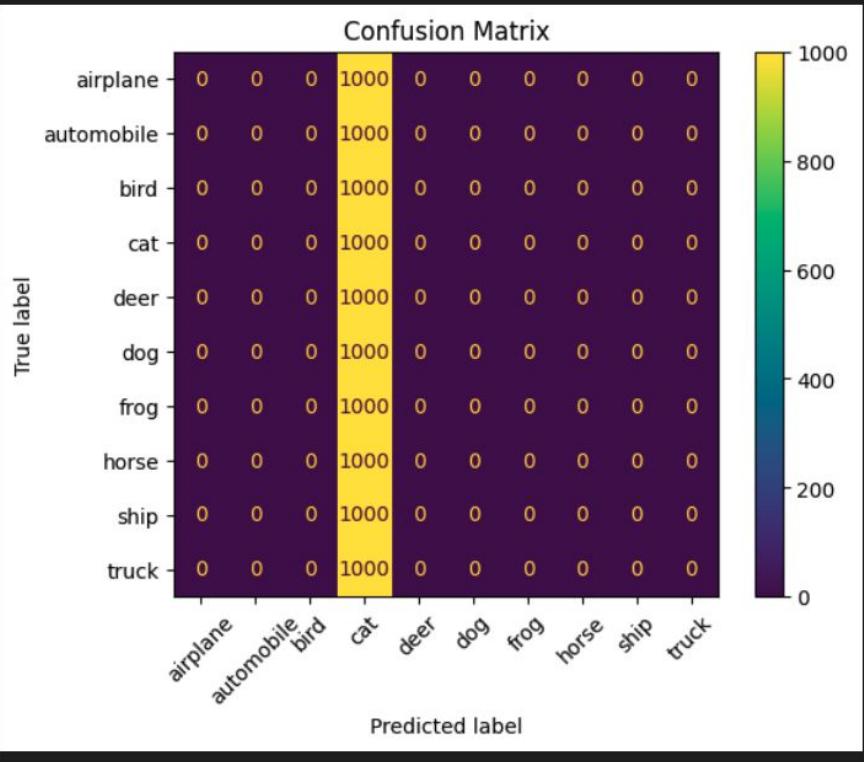
- It helps to challenge your model by manipulating the data input: think of it as giving your model the really harsh and difficult physics teacher → the model is going to do really well on the test
- Pre-trained models achieve better results, quicker, but have more overfitting.

Recommendations & Learnings

- Saving and Evaluating
 - Take some time together to think about how to save your models, and evaluations
- Create multiple notebooks to keep a better overview over your code
- Take a step back if it gets frustrating
- Pick a job that doesn't require computer vision if you hate it

What we learned?





Questions?



Yes, everything *IS* cats (not 42).

Early stopping...



train_time_sec
35802.59
3620.99
6744.72

triggered by Google Colab.

Recommendations & Learnings

- Saving and Evaluating
 - Take some time together to think about how to save your models, and evaluations
- Create multiple notebooks to keep a better overview over your code
- Take a step back if it gets frustrating
- Pick a job that doesn't require computer vision if you hate it